

detection.docx

by Muhammad Nazir

Submission date: 01-Aug-2024 11:18PM (UTC+1000)

Submission ID: 2423587036

File name: detection.docx (72.8K)

Word count: 2811

Character count: 19905

```

# Imports
1 import matplotlib.pyplot as plt
import matplotlib.patches as patches
import numpy as np
import os

import torch

```

```

from collections import Counter
from torch.utils.data import DataLoader
8 from tqdm import tqdm
import torch
import torch.nn as nn

```

```

1 config = [
    (32, 3, 1),
    (64, 3, 2),
    ["B", 1],
    (128, 3, 2),
    ["B", 2],
    (256, 3, 2),
    ["B", 8],
    (512, 3, 2),
    ["B", 8],
    (1024, 3, 2),
    ["B", 4], # To this point is Darknet-53

```

```

    (512, 1, 1),
    (1024, 3, 1),
    "S",
    (256, 1, 1),
    "U",
    (256, 1, 1),
    (512, 3, 1),
    "S",
    (128, 1, 1),
    "U",
    (128, 1, 1),
    (256, 3, 1),
    "S",
]

```

```

class CNNBlock(nn.Module):
    def __init__(self, in_channels, out_channels, bn_act=True, **kwargs):
        super(CNNBlock, self).__init__()
        self.conv = nn.Conv2d(in_channels, out_channels, bias=not bn_act,
**kwargs) # If batchnorm layer(bn_act) is true, then bias is False
1         self.bn = nn.BatchNorm2d(out_channels)
        self.leaky = nn.LeakyReLU(0.1)

```

```
self.use_bn_act = bn_act
```

```
def forward(self, x):  
    if self.use_bn_act:  
        return self.leaky(self.bn(self.conv(x)))  
    else:  
        return self.conv(x)
```

```
class ResidualBlock(nn.Module):  
    def __init__(self, channels, use_residual=True, num_repeats=1):  
        super(ResidualBlock, self).__init__()  
        self.layers = nn.ModuleList() # Like regular python list, but is  
        container for pytorch nn modules
```

```
1  
    for repeat in range(num_repeats):  
        self.layers += [  
            nn.Sequential(  
                CNNBlock(channels, channels//2, kernel_size=1),  
                CNNBlock(channels//2, channels, kernel_size=3, padding=1)  
            )  
        ]
```

```
self.use_residual = use_residual  
self.num_repeats = num_repeats
```

```
def forward(self, x):  
    for layer in self.layers:  
        if self.use_residual:  
            x = x + layer(x)  
        else:  
            x = layer(x)
```

```
return x
```

```
class ScalePrediction(nn.Module):  
    def __init__(self, in_channels, num_classes):  
        super(ScalePrediction, self).__init__()  
        self.pred = nn.Sequential(  
            CNNBlock(in_channels, 2 * in_channels, kernel_size=3,  
padding=1),  
            CNNBlock(2 * in_channels, (num_classes + 5) * 3, bn_act=False,  
kernel_size=1), # (num_classes + 5) * 3 --> (20+5) for each anchor box  
which in total is 3
```

```
1  
self.num_classes = num_classes
```

```
def forward(self, x):  
    return (  
        self.pred(x)
```

```

        .reshape(x.shape[0], 3, self.num_classes + 5, x.shape[2],
x.shape[3]) # [batch_size, anchor_boxes, prediction(25), grid_h, grid_w]
        .permute(0, 1, 3, 4, 2) # [batch_size, anchor_boxes, grid_h,
grid_w, prediction(25)]
    )

```

```

1 class YOLOv3(nn.Module):
    def __init__(self, in_channels=3, num_classes=20):
        super(YOLOv3, self).__init__()
        self.num_classes = num_classes
        self.in_channels = in_channels
        self.layers = self._create_conv_layers()

```

```

    def forward(self, x):
        1 outputs = []
        route_connections = []

```

```

        for layer in self.layers:
            if isinstance(layer, ScalePrediction):
                outputs.append(layer(x))
                continue

```

```

            x = layer(x)

```

```

            if isinstance(layer, ResidualBlock) and layer.num_repeats == 8:
                route_connections.append(x)

```

```

            elif isinstance(layer, nn.Upsample):
                x = torch.cat([x, route_connections[-1]], dim=1)
                route_connections.pop()

```

```

        return outputs

```

```

    def _create_conv_layers(self):
        layers = nn.ModuleList()
        in_channels = self.in_channels

```

```

        for module in config:
            if isinstance(module, tuple):
                out_channels, kernel_size, stride = module
                layers.append(CNNBlock(
                    in_channels,
                    out_channels,
                    kernel_size=kernel_size,
                    stride=stride,
                    padding=1 if kernel_size == 3 else 0
                ))
                in_channels = out_channels

```

```

        elif isinstance(module, list):
            num_repeats = module[1]
            layers.append(ResidualBlock(in_channels,
num_repeats=num_repeats))

```

```

        elif isinstance(module, str):
            if module == "S":
                layers += [
                    ResidualBlock(in_channels, use_residual=False,
num_repeats=1),
                    CNNBlock(in_channels, in_channels//2, kernel_size=1),
                    ScalePrediction(in_channels//2, num_classes =
self.num_classes)
                ]
            in_channels = in_channels // 2

```

```

        elif module == "U":
            layers.append(nn.Upsample(scale_factor=2))
            in_channels = in_channels * 3

```

```

        return layers
num_classes = 20
IMAGE_SIZE = 416
model = YOLOv3(num_classes=num_classes)
x = torch.randn((2, 3, IMAGE_SIZE, IMAGE_SIZE))
out = model(x)
assert model(x)[0].shape == (2, 3, IMAGE_SIZE//32, IMAGE_SIZE//32,
num_classes + 5)
assert model(x)[1].shape == (2, 3, IMAGE_SIZE//16, IMAGE_SIZE//16,
num_classes + 5)
assert model(x)[2].shape == (2, 3, IMAGE_SIZE//8, IMAGE_SIZE//8,
num_classes + 5)

```

```

import cv2
import torch

```

```

DEVICE = "cuda" if torch.cuda.is_available() else "cpu"

```

```

1
NUM_WORKERS = 4
BATCH_SIZE = 32
IMAGE_SIZE = 416
NUM_CLASSES = 20
1
LEARNING_RATE = 1e-5
NUM_EPOCHS = 80
CONF_THRESHOLD = 0.8
MAP_IOU_THRESH = 0.5
NMS_IOU_THRESH = 0.45
S = [IMAGE_SIZE // 32, IMAGE_SIZE // 16, IMAGE_SIZE // 8]

```

```
IMG_DIR = "/kaggle/input/pascalvoc-yolo/images"
LABEL_DIR = "/kaggle/input/pascalvoc-yolo/labels"
```

```
ANCHORS = [
    [(0.28, 0.22), (0.38, 0.48), (0.9, 0.78)],
    [(0.07, 0.15), (0.15, 0.11), (0.14, 0.29)],
    [(0.02, 0.03), (0.04, 0.07), (0.08, 0.06)],
]
```

```
7
PASCAL_CLASSES = [
    "aeroplane",
    "bicycle",
    "bird",
    "boat",
    "bottle",
    "bus",
    "car",
    "cat",
    "chair",
    "cow",
    "diningtable",
    "dog",
    "horse",
    "motorbike",
    "person",
    "pottedplant",
    "sheep",
    "sofa",
    "train",
    "tvmonitor"
]
```

```
def iou_width_height(boxes1, boxes2):
```

```
1
    intersection = torch.min(boxes1[..., 0], boxes2[..., 0]) *
    torch.min(
        boxes1[..., 1], boxes2[..., 1]
    )
    union = (
        boxes1[..., 0] * boxes1[..., 1] + boxes2[..., 0] * boxes2[...
1] - intersection
    )
    return intersection / union
```

```
def intersection_over_union(boxes_preds, boxes_labels,
box_format="midpoint"):
```

```
1
    if box_format == "midpoint":
```

```

box1_x1 = boxes_preds[..., 0:1] - boxes_preds[..., 2:3] / 2
box1_y1 = boxes_preds[..., 1:2] - boxes_preds[..., 3:4] / 2
box1_x2 = boxes_preds[..., 0:1] + boxes_preds[..., 2:3] / 2
box1_y2 = boxes_preds[..., 1:2] + boxes_preds[..., 3:4] / 2
box2_x1 = boxes_labels[..., 0:1] - boxes_labels[..., 2:3] / 2
box2_y1 = boxes_labels[..., 1:2] - boxes_labels[..., 3:4] / 2
box2_x2 = boxes_labels[..., 0:1] + boxes_labels[..., 2:3] / 2
box2_y2 = boxes_labels[..., 1:2] + boxes_labels[..., 3:4] / 2

```

```

if box_format == "corners":
    box1_x1 = boxes_preds[..., 0:1]
    box1_y1 = boxes_preds[..., 1:2]
    box1_x2 = boxes_preds[..., 2:3]
    box1_y2 = boxes_preds[..., 3:4]
    box2_x1 = boxes_labels[..., 0:1]
    box2_y1 = boxes_labels[..., 1:2]
    box2_x2 = boxes_labels[..., 2:3]
    box2_y2 = boxes_labels[..., 3:4]

```

```

x1 = torch.max(box1_x1, box2_x1)
y1 = torch.max(box1_y1, box2_y1)
x2 = torch.min(box1_x2, box2_x2)
y2 = torch.min(box1_y2, box2_y2)

```

```

intersection = (x2 - x1).clamp(0) * (y2 - y1).clamp(0)
box1_area = abs((box1_x2 - box1_x1) * (box1_y2 - box1_y1))
box2_area = abs((box2_x2 - box2_x1) * (box2_y2 - box2_y1))

```

```

return intersection / (box1_area + box2_area - intersection + 1e-6)

```

```

def non_max_suppression(bboxes, iou_threshold, threshold,
    box_format="corners"):
    assert type(bboxes) == list

```

```

    bboxes = [box for box in bboxes if box[1] > threshold]
    bboxes = sorted(bboxes, key=lambda x: x[1], reverse=True)
    bboxes_after_nms = []

```

```

while bboxes:
    chosen_box = bboxes.pop(0)

```

```

    bboxes = [
        box
        for box in bboxes
        if box[0] != chosen_box[0]
        or intersection_over_union(
            torch.tensor(chosen_box[2:]),
            torch.tensor(box[2:]),
            box_format=box_format,
        )
    ]

```

```

        < iou_threshold
    ]

    bboxes_after_nms.append(chosen_box)

10 return bboxes_after_nms
import numpy as np
import os
import pandas as pd
import torch

from torch.utils.data import Dataset, DataLoader
from PIL import Image, ImageFile

1 # allows PIL to load images even if they are truncated or incomplete
ImageFile.LOAD_TRUNCATED_IMAGES = True

class YOLODataset(Dataset):
    def __init__(self, csv_file, img_dir, label_dir, anchors,
                  image_size=416, S=[13,26,52], C=20, transform=None):
        self.annotations = pd.read_csv(csv_file)
        self.img_dir = img_dir
        self.label_dir = label_dir
        self.transform = transform
        self.S = S

        9 # Suppose, anchors[0] = [a,b,c], anchors[1] = [d,e,f], anchors[2] =
        [g,h,i] : Each set of anchors for each scale
        # List addition gives shape 3x3
        # Anchors per scale suggests that there are three different aspect
        ratios for each anchor position.
        1 self.anchors = torch.tensor(anchors[0] + anchors[1] + anchors[2]) #
        For all 3 scales
        self.num_anchors = self.anchors.shape[0]
        self.num_anchors_per_scale = self.num_anchors // 3

    self.C = C

    # If a cell has obj. then one anchor is responsible for outputting
    it,
    # one that's responsible is the one that has highest iou with
    ground truth box
    # but, there might be cases where there are several boxes in the
    same cell
    1 self.ignore_iou_thresh = 0.5

    def __len__(self):
        return len(self.annotations)

    def __getitem__(self, index):

```



```

label_path = os.path.join(self.label_dir,
self.annotations.iloc[index, 1])
bboxes = np.roll(np.loadtxt(fname=label_path, delimiter=" ",
min=2), 4, axis=1).tolist() # np.roll with shift 4 on axis 1: [class,
18 x, y, w, h] --> [x, y, w, h, class]

4 img_path = os.path.join(self.img_dir, self.annotations.iloc[index,
0])
image = Image.open(img_path)

if self.transform:
image = self.transform(image)

6 targets = [torch.zeros((self.num_anchors // 3, S, S, 6)) for S in
9 self.S] # 6 because objectness score, bounding box coordinates (x, y, w,
h), class label

for box in bboxes:
    """For each box in bboxes,
    2 we want to assign which anchor should be responsible and
    which cell should be responsible for all the three different
    scales prediction"""
    3 iou_anchors = iou_width_height(torch.tensor(box[2:4]),
self.anchors) # IOU from height and width
    2 anchor_indices = iou_anchors.argsort(descending=True, dim=0) #
Sorting sucht that the first is the best anchor

6 x, y, width, height, class_label = box
has_anchor = [False, False, False] # Make sure there is an anchor
for each of three scales for each bounding box

2 for anchor_idx in anchor_indices:
    scale_idx = anchor_idx // self.num_anchors_per_scale #
scale_idx is either 0,1,2: 0-->13x13, 1-->26x26, 2-->52x52
    1 anchor_on_scale = anchor_idx % self.num_anchors_per_scale # In
each scale, choosing the anchor thats either 0,1,2

    1 S = self.S[scale_idx]
    i, 1 j = int(S*y), int(S*x) # x=0.5, S=13 --> int(6.5) = 6 | i=y
cell, j=x cell
    anchor_taken = targets[scale_idx][anchor_on_scale, i, j, 0]

    if not anchor_taken and not has_anchor[scale_idx]:
        targets[scale_idx][anchor_on_scale, i, j, 0] = 1
        1 x_cell, y_cell = S*x - j, S*y - i # 6.5 - 6 = 0.5 such that
they are between [0,1]
        width_cell, height_cell = (
            width*S, # S=13, width=0.5, 6.5
            height*S
        )

```

```

    box_coordinates = torch.tensor([x_cell, y_cell, width_cell,
height_cell])

    targets[scale_idx][anchor_on_scale, i, j, 1:5] =
box_coordinates
    targets[scale_idx][anchor_on_scale, i, j, 5] =
int(class_label)
    has_anchor[scale_idx] = True

    # Even if the same grid shares another anchor having
iou>ignore iou_thresh then,
    elif not anchor_taken and iou_anchors[anchor_idx] >
self.ignore_iou_thresh:
        targets[scale_idx][anchor_on_scale, i, j, 0] = -1 # ignore
this prediction

    return image, tuple(targets)

12
import torchvision.transforms as transforms
transform = transforms.Compose([transforms.Resize((416, 416)),
transforms.ToTensor()])

1
def get_loaders(train_csv_path, test_csv_path):

    1
    train_dataset = YOLODataset(
        train_csv_path,
        transform=transform,
        S=[IMAGE_SIZE // 32, IMAGE_SIZE // 16, IMAGE_SIZE // 8],
        img_dir=IMG_DIR,
        label_dir=LABEL_DIR,
        anchors=ANCHORS,
    )
    test_dataset = YOLODataset(
        test_csv_path,
        transform=transform,
        S=[IMAGE_SIZE // 32, IMAGE_SIZE // 16, IMAGE_SIZE // 8],
        img_dir=IMG_DIR,
        label_dir=LABEL_DIR,
        anchors=ANCHORS,
    )
    train_loader = DataLoader(
        dataset=train_dataset,
        batch_size=BATCH_SIZE,
        5
        shuffle=True,
        drop_last=False,
    )
    test_loader = DataLoader(
        dataset=test_dataset,

```

```
        batch_size=BATCH_SIZE,  
        shuffle=False,  
        drop_last=False,  
    )
```

```
    return train_loader, test_loader
```

```
def mean_average_precision(  
    pred_boxes, true_boxes, iou_threshold=0.5, box_format="midpoint",  
    num_classes=4  
):
```

```
    # list storing all AP for respective classes  
    average_precisions = []
```

```
    # used for numerical stability later on  
    epsilon = 1e-6
```

```
    for c in range(num_classes):  
        detections = []  
        ground_truths = []
```

```
        # Go through all predictions and targets,  
        # and only add the ones that belong to the  
        # current class c  
        for detection in pred_boxes:  
            if detection[1] == c:  
                detections.append(detection)
```

```
        for true_box in true_boxes:  
            if true_box[1] == c:  
                ground_truths.append(true_box)
```

```
        # find the amount of bboxes for each training example  
        # Counter here finds how many ground truth bboxes we get  
        # for each training example, so let's say img 0 has 3,  
        # img 1 has 5 then we will obtain a dictionary with:  
        # amount_bboxes = {0:3, 1:5}  
        amount_bboxes = Counter([gt[0] for gt in ground_truths])
```

```
        # We then go through each key, val in this dictionary  
        # and convert to the following (w.r.t same example):  
        # ammount_bboxes = {0:torch.tensor[0,0,0],  
        1:torch.tensor[0,0,0,0,0]}  
        for key, val in amount_bboxes.items():  
            amount_bboxes[key] = torch.zeros(val)
```

```
        # sort by box probabilities which is index 2  
        detections.sort(key=lambda x: x[2], reverse=True)  
        TP = torch.zeros((len(detections)))
```

```
FP = torch.zeros((len(detections)))
total_true_bboxes = len(ground_truths)
```

```
# If none exists for this class then we can safely skip
if total_true_bboxes == 0:
    continue
```

```
for detection_idx, detection in enumerate(detections):
    # Only take out the ground_truths that have the same
    # training idx as detection
    ground_truth_img = [
        bbox for bbox in ground_truths if bbox[0] ==
detection[0]
    ]
```

```
num_gts = len(ground_truth_img)
best_iou = 0
```

```
for idx, gt in enumerate(ground_truth_img):
    iou = intersection_over_union(
        torch.tensor(detection[3:]),
        torch.tensor(gt[3:]),
        box_format=box_format,
    )
```

```
if iou > best_iou:
    best_iou = iou
    best_gt_idx = idx
```

```
if best_iou > iou_threshold:
    # only detect ground truth detection once
    if amount_bboxes[detection[0]][best_gt_idx] == 0:
        # true positive and add this bounding box to seen
        TP[detection_idx] = 1
        amount_bboxes[detection[0]][best_gt_idx] = 1
    else:
        FP[detection_idx] = 1
```

```
# if IOU is lower then the detection is a false positive
else:
    FP[detection_idx] = 1
```

```
TP_cumsum = torch.cumsum(TP, dim=0)
FP_cumsum = torch.cumsum(FP, dim=0)
recalls = TP_cumsum / (total_true_bboxes + epsilon)
precisions = TP_cumsum / (TP_cumsum + FP_cumsum + epsilon)
precisions = torch.cat((torch.tensor([1]), precisions))
recalls = torch.cat((torch.tensor([0]), recalls))
# torch.trapz for numerical integration
average_precisions.append(torch.trapz(precisions, recalls))
```

```

return sum(average_precisions) / len(average_precisions)

def get_evaluation_bboxes(
    loader,
    model,
    iou_threshold,
    anchors,
    threshold,
    box_format="midpoint",
    device="cuda" if torch.cuda.is_available() else "cpu",
):
    # make sure model is in eval before get bboxes
    model.eval()
    train_idx = 0
    all_pred_boxes = []
    all_true_boxes = []
    for batch_idx, (x, labels) in enumerate(loader):
        x = x.float().to(device)

        with torch.no_grad():
            predictions = model(x)

            batch_size = x.shape[0]
            bboxes = [[] for _ in range(batch_size)]
            for i in range(3):
                S = predictions[i].shape[2] # grid cell size for each
                anchor = torch.tensor([*anchors[i]]).to(device) * S #
                anchor for each grid, prediction type
                boxes_scale_i = cells_to_bboxes( # get bboxes for each
                    image in the batch
                    predictions[i], anchor, S=S, is_preds=True
                )
                for idx, (box) in enumerate(boxes_scale_i): # for each
                    image, append the bbox to corr. bboxes[idx]
                    bboxes[idx] += box

            # we just want one bbox for each label, not one for each scale
            true_bboxes = cells_to_bboxes(
                labels[2], anchor, S=S, is_preds=False
            )

            for idx in range(batch_size):
                nms_boxes = non_max_suppression(
                    bboxes[idx],
                    iou_threshold=iou_threshold,
                    threshold=threshold,
                    box_format=box_format,

```

```

    )

    for nms_box in nms_boxes:
        all_pred_boxes.append([train_idx] + nms_box)

    for box in true_bboxes[idx]:
        if box[1] > threshold:
            all_true_boxes.append([train_idx] + box)

    train_idx += 1

model.train()
return all_pred_boxes, all_true_boxes

def cells_to_bboxes(predictions, anchors, S, is_preds=True):
    1
    BATCH_SIZE = predictions.shape[0]
    num_anchors = len(anchors)
    box_predictions = predictions[..., 1:5]
    if is_preds:
        anchors = anchors.reshape(1, len(anchors), 1, 1, 2)
        box_predictions[..., 0:2] = torch.sigmoid(box_predictions[...,
0:2])
        box_predictions[..., 2:] = torch.exp(box_predictions[..., 2:])
    * anchors
    scores = torch.sigmoid(predictions[..., 0:1])
    best_class = torch.argmax(predictions[..., 5:], dim=-
1).unsqueeze(-1)
    else:
        scores = predictions[..., 0:1]
        best_class = predictions[..., 5:6]

    cell_indices = (
        torch.arange(S)
        .repeat(predictions.shape[0], 3, S, 1)
        .unsqueeze(-1)
        .to(predictions.device)
    )
    x = 1 / S * (box_predictions[..., 0:1] + cell_indices)
    y = 1 / S * (box_predictions[..., 1:2] + cell_indices.permute(0, 1,
3, 2, 4))
    w_h = 1 / S * box_predictions[..., 2:4]
    converted_bboxes = torch.cat((best_class, scores, x, y, w_h), dim=-
1).reshape(BATCH_SIZE, num_anchors * S * S, 6)
    return converted_bboxes.tolist()

    2
class YoloLoss(nn.Module):
    def __init__(self):
        super(YoloLoss, self).__init__()

```

```

self.mse = nn.MSELoss() # For bounding box loss
self.bce = nn.BCEWithLogitsLoss() # For multi-label prediction:
Binary cross entropy
1 self.entropy = nn.CrossEntropyLoss() # For classification
self.sigmoid = nn.Sigmoid()

1 Constants for significance of obj, or no obj.
self.lambda_class = 1
self.lambda_noobj = 10
self.lambda_obj = 1
self.lambda_box = 10

def 2 forward(self, predictions, target, anchors):
    obj = target[..., 0] == 1
    noobj = target[..., 0] == 0

    # No object Loss
    #####
    no_object_loss = self.bce(
        (predictions[..., 0:1][noobj]), (target[..., 0:1][noobj])
    )

    # Object Loss
    #####
    anchors = anchors.reshape(1,3,1,1,2) # Anchors initial shape 3x2 -->
3 anchor boxes each of certain hwx (2)

    # box_preds = [..., sigmoid(x), sigmoid(y), [p_w * exp(t_w)], [p_h
    * exp(t_h)], ...]
    19
    1 box_preds = torch.cat([self.sigmoid(predictions[..., 1:3]),
    torch.exp(predictions[..., 3:5]) * anchors], dim=-1)

    1 iou between predicted box and target box
    ious = intersection_over_union(box_preds[obj], target[...,
    1:5][obj]).detach()

    object_loss = self.bce(
        (predictions[..., 0:1][obj]), (ious * target[..., 0:1][obj]) #
    target * iou because only intersected part object loss calc
    )

    8
    # Box Coordinate Loss
    #####
    predictions[..., 1:3] = self.sigmoid(predictions[..., 1:3]) # x, y
    to b 2 between [0,1]
    target[..., 3:5] = torch.log(
        (1e-6 + target[..., 3:5] / anchors)
    ) # Exponential of hwx (taking log because opp. of exp)

```

```
1 box_loss = self.mse(predictions[..., 1:5][obj], target[..., 1:5][obj])
```

```
# Class Loss
#####
class_loss = self.entropy(
    (predictions[..., 5:][obj]), (target[..., 5][obj].long())
)
```

```
return(
    self.lambda_box * box_loss
    + self.lambda_obj * object_loss
    + self.lambda_noobj * no_object_loss
    + self.lambda_class * class_loss
)
```

```
1 def plot_image(image, boxes):
    """Plots predicted bounding boxes on the image"""
    cmap = plt.get_cmap("tab20b")
    class_labels = PASCAL_CLASSES
    colors = [cmap(i) for i in np.linspace(0, 1, len(class_labels))]
    im = np.array(image)
    height, width, _ = im.shape
```

```
# Create figure and axes
fig, ax = plt.subplots(1)
# Display the image
ax.imshow(im)
```

```
# box[0] is x midpoint, box[2] is width
# box[1] is y midpoint, box[3] is height
```

```
# Create a Rectangle patch
for box in boxes:
    assert len(box) == 6, "box should contain class pred,
confidence, x, y, width, height"
    class_pred = box[0]
    box = box[2:]
    upper_left_x = box[0] - box[2] / 2
    upper_left_y = box[1] - box[3] / 2
    rect = patches.Rectangle(
        (upper_left_x * width, upper_left_y * height),
        box[2] * width,
        box[3] * height,
        linewidth=2,
        edgecolor=colors[int(class_pred)],
        facecolor="none",
    )
    # Add the patch to the Axes
```



```

        ax.add_patch(rect)
        plt.text(
            upper_left_x * width,
            upper_left_y * height,
            s=class_labels[int(class_pred)],
            color="white",
            verticalalignment="top",
            bbox={"color": colors[int(class_pred)], "pad": 0},
        )

```

```
plt.show()
```

```
# Save test loader to a file
```

```
torch.save(test_loader, '/kaggle/working/test_loader.pth')
```

```
15 import torch.optim as optim
```

```
from tqdm import tqdm
```

```
import time
```

```
17 history_loss = [] # To plot the epoch vs. loss
```

```
for epoch in tqdm(range(NUM_EPOCHS), desc="Epochs"):
    model.train()
```

```
    losses = []
```

```
    start_time = time.time() # Start time of the epoch
```

```
1 for batch_idx, (x,y) in enumerate(train_loader):
    x = x.to(DEVICE)
    y0, y1, y2 = (y[0].to(DEVICE),
                  y[1].to(DEVICE),
                  y[2].to(DEVICE))
```

```
# context manager is used in PyTorch to automatically handle mixed-precision computations on CUDA-enabled GPUs
```

```
1 with torch.cuda.amp.autocast():
    out = model(x)
    loss = (
        loss_fn(out[0], y0, scaled_anchors[0])
        + loss_fn(out[1], y1, scaled_anchors[1])
        + loss_fn(out[2], y2, scaled_anchors[2])
    )
```

```
    losses.append(loss.item())
```

```
    optimizer.zero_grad()
```

```
    scaler.scale(loss).backward()
```

```
    scaler.step(optimizer)
```

```

        scaler.update()

    end_time = time.time() # End time of the epoch
    epoch_duration = end_time - start_time # Duration of the epoch

    history_loss.append(sum(losses)/len(losses))

    if (epoch+1) % 10 == 0:
        # Print the epoch duration
        tqdm.write(f"Epoch {epoch+1} completed in {epoch_duration:.2f}
seconds")

    # Print the loss and accuracy for training and validation data
    print(f"Epoch [{epoch+1}/{NUM_EPOCHS}], "
          f"Loss: {sum(losses)/len(losses):.4f}")

    # save the model after every 10 epoch
    torch.save(model.state_dict(),
f'/kaggle/working/Yolov3_epoch{epoch+1}.pth')

import matplotlib.pyplot as plt
epochs = range(1, len(history_loss)+1)

# Instantiate the model
model = YOLOv3(num_classes=NUM_CLASSES).to(DEVICE)

# Compile the model
optimizer = torch.optim.Adam(
    model.parameters(), lr=LEARNING_RATE

1
loss_fn = YoloLoss()

# Scaler
scaler = torch.cuda.amp.GradScaler()

# Train-Test Loader
train_loader, test_loader = get_loaders(
    train_csv_path='/kaggle/input/pascalvoc-yolo/test.csv',
    test_csv_path='/kaggle/input/pascalvoc-yolo/test.csv'
)

# Anchors
2
scaled_anchors = (
    torch.tensor(ANCHORS)*
    torch.tensor([13,26,52]).unsqueeze(1).unsqueeze(1).repeat(1,3,2)
).to(DEVICE)

1
model.eval()
x, y = next(iter(test_loader))
x = x.float().to(DEVICE)

```

```

with torch.no_grad():
    out = model(x)
    bboxes = [[] for _ in range(x.shape[0])]
    batch_size, A, S, _, _ = out[0].shape
    anchor = torch.tensor([*ANCHORS[0]]).to(DEVICE) * S
    boxes_scale_i = cells_to_bboxes(
        out[0], anchor, S=S, is_preds=True
    )
    for idx, (box) in enumerate(boxes_scale_i):
        bboxes[idx] += box

```

```

    for i in range(batch_size):
        nms_boxes = non_max_suppression(
            bboxes[i], iou_threshold=0.5, threshold=0.6,
            box_format="midpoint",
        )
        plot_image(x[i].permute(1,2,0).detach().cpu(), nms_boxes)
# Load the model
model = YOLOv3(num_classes=NUM_CLASSES)
model_path = "/kaggle/input/80-epoch-yolov3-model/Yolov3_epoch80.pth"
state_dict = torch.load(model_path)
model.load_state_dict(state_dict)
model = model.to(DEVICE)

```

```

# Testing
losses = []

```

```

11 with torch.no_grad():
    model.eval()

```

```

    for batch_idx, (x,y) in enumerate(test_loader):
        x = x.to(DEVICE)
        y0, y1, y2 = (y[0].to(DEVICE),
                      y[1].to(DEVICE),
                      y[2].to(DEVICE))

```

```

        out = model(x)
        loss = (
            loss_fn(out[0], y0, scaled_anchors[0])
            + loss_fn(out[1], y1, scaled_anchors[1])
            + loss_fn(out[2], y2, scaled_anchors[2])
        )

```

```

        losses.append(loss.item())

```

```

print(f"Loss: {sum(losses)/len(losses):.4f}")

```

```

1 pred_boxes, true_boxes = get_evaluation_bboxes(
    test_loader,

```

```
        model,  
        iou_threshold=NMS_IOU_THRESH,  
        anchors=ANCHORS,  
        threshold=CONF_THRESHOLD,  
    )  
  
    mapval = mean_average_precision(  
        pred_boxes,  
        true_boxes,  
        iou_threshold=MAP_IOU_THRESH,  
        box_format="midpoint",  
        num_classes=NUM_CLASSES,  
    )  
    print(f"MAP: {mapval.item()}")
```

detection.docx

ORIGINALITY REPORT

81%

SIMILARITY INDEX

81%

INTERNET SOURCES

13%

PUBLICATIONS

28%

STUDENT PAPERS

PRIMARY SOURCES

1

huggingface.co

Internet Source

71%

2

origin.geeksforgeeks.org

Internet Source

3%

3

discuss.pytorch.org

Internet Source

1%

4

Submitted to University of East London

Student Paper

1%

5

Submitted to University of Sydney

Student Paper

1%

6

u-n-joe.tistory.com

Internet Source

1%

7

git.trustie.net

Internet Source

1%

8

hdl.handle.net

Internet Source

1%

9

mdpi-res.com

Internet Source

<1%

10	Submitted to City University Student Paper	<1 %
11	Submitted to University of North Texas Student Paper	<1 %
12	ctkim.tistory.com Internet Source	<1 %
13	Eklas Hossain. "Machine Learning Crash Course for Engineers", Springer Science and Business Media LLC, 2024 Publication	<1 %
14	deep-learning-study.tistory.com Internet Source	<1 %
15	kypseli.ouc.ac.cy Internet Source	<1 %
16	openi.pcl.ac.cn Internet Source	<1 %
17	raschka-research-group.github.io Internet Source	<1 %
18	velog.io Internet Source	<1 %
19	Shancheng Fang, Hongtao Xie, Zhineng Chen, Shiai Zhu, Xiaoyan Gu, Xingyu Gao. "Detecting Uyghur text in complex background images with convolutional neural network", Multimedia Tools and Applications, 2017	<1 %

20

www.codetd.com

Internet Source

<1 %

21

Adnan, Tamim. "Classification and Pixel-Level Segmentation of Asphalt Pavement Cracks Using Convolutional Neural Networks", The University of North Carolina at Charlotte, 2023

Publication

<1 %

Exclude quotes Off

Exclude matches Off

Exclude bibliography On